


Beyond the Report: When Foresight Becomes Infrastructure

*But if we can observe change continuously, why should
foresight still behave like a publication?*



A report begins ageing the moment it is published.

STATIC OUTPUT

The document stays exactly as published.

Foresight has always had a peculiar relationship with time. We study change in order to understand what might come next, yet the result of that work is often frozen into a document that begins ageing the moment it is published. Signals keep emerging, technologies converge, assumptions weaken, new evidence contradicts old conclusions, but the report remains exactly as it was on publication day.

For a long time, this was simply a limitation of the medium. Research was expensive, information scarce enough to be manageable, and updating an analysis meant essentially starting the process again. AI changes that equation. We can now scan, classify, connect and interrogate volumes of information that would have been impossible for a human research team to process continuously.

CONTINUOUS CHANGE

Signals, evidence and assumptions keep moving.

Over the past months, we have been exploring this question by designing and prototyping an AI-native foresight system. The work was developed in a professional context, so the architecture, implementation and outputs remain proprietary. What is more interesting to share anyway is what happened to the underlying idea of foresight once we tried to turn it into something that actually had to work.

Because building the system exposed a much harder question than how to automate research: what exactly are we trying to automate?

**But if we can observe change continuously,
why should foresight still behave like a publication?**

Finding information is not the same as understanding change.

At first glance, AI and foresight appear almost perfectly matched. Foresight requires scanning large amounts of information for indications of change, while AI has become remarkably good at processing enormous volumes of unstructured information. Give a system access to research papers, patents, investment activity, company announcements, regulatory developments and technological news, and the obvious promise is that it will find signals humans might miss.

And it probably will.

But finding information is not the same as understanding change.

A new patent is not a trend. Neither is a startup announcement, a research breakthrough or a surge in media attention. Even thousands of signals do not automatically add up to foresight. Their meaning emerges from the relationships between them: technologies beginning to converge, capabilities migrating from one domain into another, previously isolated experiments becoming commercially viable, regulation altering adoption trajectories, or multiple weak developments suddenly reinforcing one another.

This is where the apparent simplicity of AI-powered foresight begins to disappear. The challenge is no longer scarcity of information but abundance without meaning. If yesterday's problem was finding enough evidence, tomorrow's may be deciding which evidence deserves our attention.

An AI-native foresight system therefore cannot simply become an extremely efficient search engine. It has to help us make sense of change.

01

Information

research / patents /
investment / regulation

02

Relationships

convergence / migration /
reinforcement

03

Meaning

an interpretation of change

The challenge is no longer scarcity of information but **abundance without meaning.**

A trend is closer to a hypothesis than a static fact.

HYPOTHESIS

AI agents become increasingly autonomous economic actors

SUPPORTING EVIDENCE

capabilities / protocols / investment / early deployments

CHALLENGING EVIDENCE

security failures / regulation / constraints / adoption

Building something forces definitions that conceptual discussions allow us to avoid.

We speak about signals, trends, megatrends and emerging technologies as if their boundaries were self-evident. They are not. At what point does a collection of signals become a trend? How much evidence is enough? How do we distinguish an accelerating structural shift from a temporary explosion of hype? And perhaps most importantly, who decides?

The more we worked on the problem, the less useful it seemed to treat a trend as a static fact. A trend is closer to a hypothesis: an interpretation of change supported by evidence that can strengthen, weaken or evolve over time.

Take the idea that AI agents could develop from digital assistants into increasingly autonomous economic actors. Today, we might find supporting evidence in new technical capabilities, emerging protocols, investment, early deployments and changes in how companies design digital processes. Tomorrow, security failures, regulation, technical constraints or disappointing adoption could challenge the same hypothesis. Neither state represents a definitive answer. What changes is the weight and structure of the evidence around it.

This creates a very different model of foresight. Instead of periodically declaring that a trend exists, we can begin asking how our confidence in it evolves. What new evidence appeared? What contradicted our assumptions? Which adjacent developments strengthened the hypothesis? Where is uncertainty increasing rather than decreasing?

The future stops looking like a catalogue of trends and begins to resemble a landscape of competing interpretations.

CONFIDENCE IS NOT A VERDICT. IT CHANGES AS THE EVIDENCE CHANGES.

The same development can mean four different things.

But matter to whom?

Once hundreds of developments can be identified, the next temptation is to rank them.

Surely AI can calculate which trends matter most.

A technological development may be existential for one industry and almost irrelevant to another. Something with enormous transformative potential may still be too immature to influence decisions today, while a mature technology may be strategically uninteresting precisely because it is already accessible to everyone. A trend can accelerate rapidly while remaining commercially irrelevant. Another may progress quietly for years before suddenly becoming critical when combined with something else.

The more dimensions we introduced - strength of evidence, acceleration, maturity, impact, time horizon, uncertainty, relevance - the less convincing the idea of one universal score became.

Context refuses to disappear.

A breakthrough in robotics may represent an automation opportunity for a manufacturer, a new category of workplace risk for an insurer, an infrastructure opportunity for a technology provider, and a regulatory challenge for government. The technology is the same; the strategic meaning is not.

This may ultimately be one of the most consequential shifts enabled by AI. Not the ability to produce more foresight, but the ability to interpret the same evolving evidence through many different lenses.

ONE TECHNOLOGY / DIFFERENT STRATEGIC MEANING

Manufacturer

automation opportunity

Insurer

workplace risk

Tech provider

infrastructure opportunity

Government


regulatory challenge

Traditional reports have to generalise. Intelligence systems do not necessarily have to.

The goal is not to eliminate uncertainty.

It is to **understand uncertainty better.**

TRACE BACK THROUGH THE REASONING



There is an appealing but dangerous version of this future in which enough data, sufficiently powerful models and a sophisticated scoring system finally allow the machine to tell us what comes next.

I don't think that is foresight.

Foresight operates precisely where certainty disappears. Evidence is incomplete, causality ambiguous and multiple interpretations remain plausible. The objective cannot be to eliminate uncertainty by generating increasingly confident answers. It should be to understand uncertainty better.

That makes explainability more important, not less, as AI becomes embedded in the research process. If a system suggests that a development is becoming strategically relevant, we should be able to move backwards through its reasoning: from implication to trend, from trend to supporting and contradicting evidence, and from evidence to the original signals from which the interpretation emerged.

Otherwise we risk replacing one form of opacity with another. Instead of analysts producing conclusions whose assumptions remain invisible, we get algorithms producing conclusions whose assumptions are even harder to interrogate.

AI can generate interpretations at extraordinary speed. That makes the ability to challenge those interpretations more valuable than ever.

Perhaps the future analyst will spend less time searching for information and more time disagreeing intelligently with machines.

A prototype forces an idea to collide with reality.



This is also where the project changed something in the way I think about research itself.

We could have explored the concept of an AI-native foresight system on paper. Many of its principles sound convincing at the conceptual level:

continuous scanning, dynamic trend detection, contextual scoring, evidence trails, personalised views. It is only when these ideas have to become operational that their hidden assumptions become visible.

What exactly is a signal? When should two pieces of evidence be connected? How long does an old signal remain relevant? When does new evidence justify changing an assessment? How should contradictory sources affect confidence? Can the system distinguish genuine acceleration from an increase in media attention? How much interpretation should be delegated to AI before human judgement becomes ceremonial?

These questions did not appear because the original idea was wrong. They appeared because something had to work.

That is what makes prototypes so valuable as research instruments. A prototype is not merely an unfinished product; it is a way of forcing an idea to collide with reality. It turns assumptions into decisions, and decisions create things that can fail.

For someone who started her career as an engineer, I should not find this surprising. There is a particular kind of knowledge that only appears when an abstract idea becomes something you can interact with, test and break. Somewhere along the way, I had started treating foresight primarily as a way of thinking about the future. Building this prototype reminded me that we can also learn about futures by making fragments of them tangible.

*Failure, in this context, is not an inconvenience. **It is data.***

Perhaps we have confused the output with the capability.

I don't believe foresight reports are disappearing. A good report does something an intelligence system cannot easily replace: it creates a coherent argument. It deliberately reduces complexity, gives people a common frame of reference and creates a moment in which an organisation can stop and consider what may be changing around it.

A report could instead become one expression of a much larger, continuously evolving system of signals, evidence, hypotheses, relationships, uncertainties and contextual interpretations. A dashboard might be another expression. A conversation with an AI system could be another. A customer-specific perspective could be another. And when a development appears important enough, perhaps the next output should not be another slide at all, but an experiment.

This is where foresight becomes particularly interesting to me.

If we can move continuously from observing change, to interpreting it, to forming hypotheses, to building something that allows us to test those hypotheses, the boundary between foresight and applied research starts to blur.

ONE CAPABILITY / MANY EXPRESSIONS

**CONTINUOUSLY EVOLVING
FORESIGHT CAPABILITY**

REPORT

DASHBOARD

CONVERSATION

EXPERIMENT

A report can be one expression of a living capability - not the capability itself.

The system can keep evidence, hypotheses, relationships, uncertainties and contextual interpretations in motion between publication moments.

What can we build today to find out?



WE STOP ASKING ONLY:

What might happen?

WE ALSO BEGIN ASKING:

What would have to be true for this future to happen, and what can we build today to find out?

AI may make foresight faster. It may make it broader and considerably more personalised. But I suspect its deeper impact will be forcing us to reconsider what foresight actually is.

Perhaps the future of foresight is not a better report about change.

Perhaps it is infrastructure for continuously exploring it.